

ASSIGNMENT - 2

Downloaded the given dataset

```
import pandas as pd
import seaborn as sns
import numpy as np
from matplotlib import pyplot as plt
%matplotlib inline
```

uploaded the given dataset

```
url = 'https://drive.google.com/file/d/160K6XcuYDyRBPgJ-
JsqThkyFoJhCv0Wy/view?usp=sharing'
path = 'https://drive.google.com/uc?
export=download&id='+url.split('/')[ -2]
df= pd.read_csv(path)
```

df

	RowNumber	CustomerId	Surname	CreditScore	Geography	Gender
Age \						
0	1	15634602	Hargrave	619	France	Female
42						
1	2	15647311	Hill	608	Spain	Female
41						
2	3	15619304	Onio	502	France	Female
42						
3	4	15701354	Boni	699	France	Female
39						
4	5	15737888	Mitchell	850	Spain	Female
43						
...
...						
9995	9996	15606229	Obijiaku	771	France	Male
39						
9996	9997	15569892	Johnstone	516	France	Male
35						
9997	9998	15584532	Liu	709	France	Female
36						
9998	9999	15682355	Sabbatini	772	Germany	Male
42						
9999	10000	15628319	Walker	792	France	Female
28						

	Tenure	Balance	NumOfProducts	HasCrCard	IsActiveMember	\
0	2	0.00	1	1		1
1	1	83807.86	1	0		1
2	8	159660.80	3	1		0
3	1	0.00	2	0		0
4	2	125510.82	1	1		1

...
9995	5	0.00	2	1	0
9996	10	57369.61	1	1	1
9997	7	0.00	1	0	1
9998	3	75075.31	2	1	0
9999	4	130142.79	1	1	0

	EstimatedSalary	Exited
0	101348.88	1
1	112542.58	0
2	113931.57	1
3	93826.63	0
4	79084.10	0

...
9995	96270.64	0
9996	101699.77	0
9997	42085.58	1
9998	92888.52	1
9999	38190.78	0


[10000 rows x 14 columns]

perform below visualizations.

1.Univariate Analysis

```
sns.displot(df.Age)
```


```
<seaborn.axisgrid.FacetGrid at 0x7f5cf9021b10>
```



2.Bi-Variate Analysis

`df.plot.line()`

`<matplotlib.axes._subplots.AxesSubplot at 0x7f5cf8e8d3d0>`



3.MultiVariate Analysis

`pip install seaborn`

```
Looking in indexes: https://pypi.org/simple, https://us-
python.pkg.dev/colab-wheels/public/simple/
Requirement already satisfied: seaborn in
/usr/local/lib/python3.7/dist-packages (0.11.2)
Requirement already satisfied: matplotlib>=2.2 in
/usr/local/lib/python3.7/dist-packages (from seaborn) (3.2.2)
Requirement already satisfied: scipy>=1.0 in
/usr/local/lib/python3.7/dist-packages (from seaborn) (1.7.3)
Requirement already satisfied: pandas>=0.23 in
/usr/local/lib/python3.7/dist-packages (from seaborn) (1.3.5)
Requirement already satisfied: numpy>=1.15 in
/usr/local/lib/python3.7/dist-packages (from seaborn) (1.21.6)
Requirement already satisfied: kiwisolver>=1.0.1 in
/usr/local/lib/python3.7/dist-packages (from matplotlib>=2.2->seaborn)
(1.4.4)
Requirement already satisfied: cycler>=0.10 in
/usr/local/lib/python3.7/dist-packages (from matplotlib>=2.2->seaborn)
(0.11.0)
Requirement already satisfied: pyparsing!=2.0.4,!=2.1.2,!
=2.1.6,>=2.0.1 in /usr/local/lib/python3.7/dist-packages (from
matplotlib>=2.2->seaborn) (3.0.9)
Requirement already satisfied: python-dateutil>=2.1 in
/usr/local/lib/python3.7/dist-packages (from matplotlib>=2.2->seaborn)
(2.8.2)
```

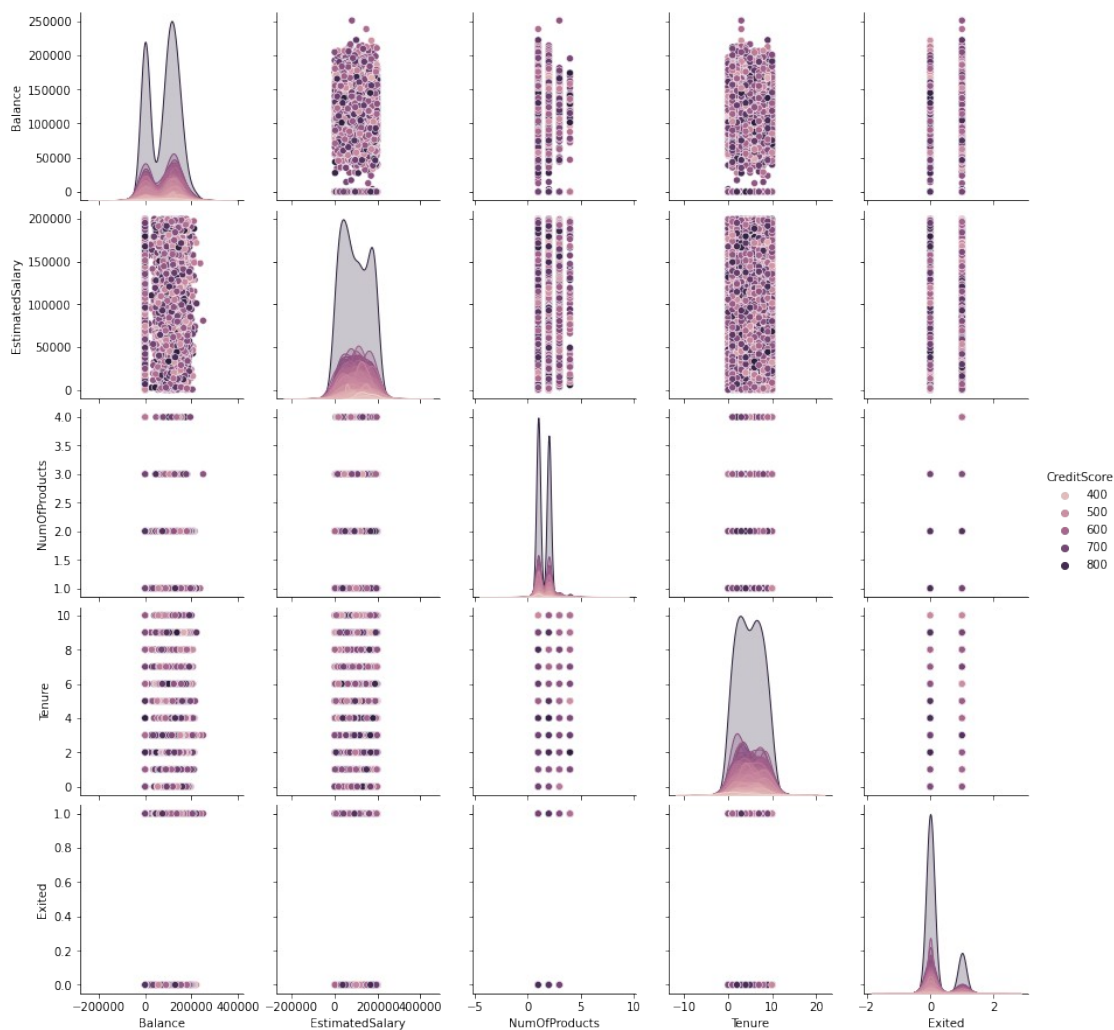
Requirement already satisfied: typing-extensions in
/usr/local/lib/python3.7/dist-packages (from kiwisolver>=1.0.1-
>matplotlib>=2.2->seaborn) (4.1.1)
Requirement already satisfied: pytz>=2017.3 in
/usr/local/lib/python3.7/dist-packages (from pandas>=0.23->seaborn)
(2022.2.1)
Requirement already satisfied: six>=1.5 in
/usr/local/lib/python3.7/dist-packages (from python-dateutil>=2.1-
>matplotlib>=2.2->seaborn) (1.15.0)

```
import seaborn as sns
```

```
plt.figure(figsize=(4,4))  
sns.pairplot(data=df[["Balance", "CreditScore", "EstimatedSalary", "NumOf  
Products", "Tenure", "Exited"]], hue="CreditScore")
```

<seaborn.axisgrid.PairGrid at 0x7f5cf301c710>

<Figure size 288x288 with 0 Axes>



Perform descriptive statistics on the dataset

```
df.describe()
```

	RowNumber	CustomerId	CreditScore	Age
Tenure \				
count	10000.000000	1.000000e+04	10000.000000	10000.000000
mean	5000.500000	1.569094e+07	650.528800	38.921800
std	2886.89568	7.193619e+04	96.653299	10.487806
min	1.000000	1.556570e+07	350.000000	18.000000
25%	2500.750000	1.562853e+07	584.000000	32.000000
50%	5000.500000	1.569074e+07	652.000000	37.000000
75%	7500.250000	1.575323e+07	718.000000	44.000000
max	10000.000000	1.581569e+07	850.000000	92.000000

	Balance	NumOfProducts	HasCrCard	IsActiveMember \
count	10000.000000	10000.000000	10000.000000	10000.000000
mean	76485.889288	1.530200	0.70550	0.515100
std	62397.405202	0.581654	0.45584	0.499797
min	0.000000	1.000000	0.00000	0.000000
25%	0.000000	1.000000	0.00000	0.000000
50%	97198.540000	1.000000	1.00000	1.000000
75%	127644.240000	2.000000	1.00000	1.000000
max	250898.090000	4.000000	1.00000	1.000000

	EstimatedSalary	Exited
count	10000.000000	10000.000000
mean	100090.239881	0.203700
std	57510.492818	0.402769
min	11.580000	0.000000
25%	51002.110000	0.000000
50%	100193.915000	0.000000
75%	149388.247500	0.000000
max	199992.480000	1.000000

Handle the missing values

```
url = 'https://drive.google.com/file/d/160K6XcuYDyRBPGj-JsqThkyFoJhCv0Wy/view?usp=sharing'
path = 'https://drive.google.com/uc?export=download&id='+url.split('/')[ -2]
```

```
df= pd.read_csv(path)
pd.isnull(df["Age"])

0      False
1      False
2      False
3      False
4      False
...
9995   False
9996   False
9997   False
9998   False
9999   False
Name: Age, Length: 10000, dtype: bool
```

Find the outliers and replace the outliers

```
df["Age"]=np.where(df["Age"]>10,np.median(df["Age"]))
df["Age"]

0      <function median at 0x7f5d15042b00>
1      <function median at 0x7f5d15042b00>
2      <function median at 0x7f5d15042b00>
3      <function median at 0x7f5d15042b00>
4      <function median at 0x7f5d15042b00>
...
9995   <function median at 0x7f5d15042b00>
9996   <function median at 0x7f5d15042b00>
9997   <function median at 0x7f5d15042b00>
9998   <function median at 0x7f5d15042b00>
9999   <function median at 0x7f5d15042b00>
Name: Age, Length: 10000, dtype: object
```

Check for categorical columns and perform encoding.

```
from sklearn.preprocessing import LabelEncoder
df['Gender'].unique()

array(['Female', 'Male'], dtype=object)

df['Gender'].value_counts()

2736    1
4076    1
8015    1
4068    1
1311    1
...
1313    1
```

```

5472    1
3785    1
4225    1
2497    1
Name: Gender, Length: 10000, dtype: int64

```

```

encoding=LabelEncoder()
df["Gender"]=encoding.fit_transform(df.iloc[:,1].values)
df

```

	RowNumber	CustomerId	Surname	CreditScore	Geography	Gender
\0	1	15634602	Hargrave	619	France	2736
1	2	15647311	Hill	608	Spain	3258
2	3	15619304	Onio	502	France	2104
3	4	15701354	Boni	699	France	5435
4	5	15737888	Mitchell	850	Spain	6899
...
9995	9996	15606229	Obijiaku	771	France	1599
9996	9997	15569892	Johnstone	516	France	161
9997	9998	15584532	Liu	709	France	717
9998	9999	15682355	Sabbatini	772	Germany	4656
9999	10000	15628319	Walker	792	France	2497

	Age	Tenure	Balance
NumOfProducts \0	<function median at 0x7f5d15042b00>	2	0.00
1	<function median at 0x7f5d15042b00>	1	83807.86
1	<function median at 0x7f5d15042b00>	8	159660.80
2	<function median at 0x7f5d15042b00>	1	0.00
3	<function median at 0x7f5d15042b00>	2	125510.82
3	<function median at 0x7f5d15042b00>		
2			
4			
1			
...
...			


```

9995 <function median at 0x7f5d15042b00>      5      0.00
2
9996 <function median at 0x7f5d15042b00>     10    57369.61
1
9997 <function median at 0x7f5d15042b00>      7      0.00
1
9998 <function median at 0x7f5d15042b00>      3    75075.31
2
9999 <function median at 0x7f5d15042b00>      4   130142.79
1

```

	HasCrCard	IsActiveMember	EstimatedSalary	Exited
0	1	1	101348.88	1
1	0	1	112542.58	0
2	1	0	113931.57	1
3	0	0	93826.63	0
4	1	1	79084.10	0
...
9995	1	0	96270.64	0
9996	1	1	101699.77	0
9997	0	1	42085.58	1
9998	1	0	92888.52	1
9999	1	0	38190.78	0

[10000 rows x 14 columns]

Split the data into dependent and independent variables

```

x=df.iloc[:, :-2].values
print(x)

[[1 15634602 'Hargrave' ... 1 1 1]
 [2 15647311 'Hill' ... 1 0 1]
 [3 15619304 'Onio' ... 3 1 0]
 ...
 [9998 15584532 'Liu' ... 1 0 1]
 [9999 15682355 'Sabbatini' ... 2 1 0]
 [10000 15628319 'Walker' ... 1 1 0]]

y=df.iloc[:, -1].values
print(y)

[1 0 1 ... 1 1 0]

```

Scale the independent variables

```

import pandas as pd
from sklearn.preprocessing import MinMaxScaler
scaler=MinMaxScaler()

```

```
df[["RowNumber"]] = scaler.fit_transform(df[["RowNumber"]])
print(df)
```

	RowNumber	CustomerId	Surname	CreditScore	Geography	Gender
\ 0	0.0000	15634602	Hargrave	619	France	2736
1	0.0001	15647311	Hill	608	Spain	3258
2	0.0002	15619304	Onio	502	France	2104
3	0.0003	15701354	Boni	699	France	5435
4	0.0004	15737888	Mitchell	850	Spain	6899
...
9995	0.9996	15606229	Obijiaku	771	France	1599
9996	0.9997	15569892	Johnstone	516	France	161
9997	0.9998	15584532	Liu	709	France	717
9998	0.9999	15682355	Sabbatini	772	Germany	4656
9999	1.0000	15628319	Walker	792	France	2497

	Age	Tenure	Balance
NumOfProducts \ 0	<function median at 0x7f5d15042b00>	2	0.00
1	<function median at 0x7f5d15042b00>	1	83807.86
1	<function median at 0x7f5d15042b00>	8	159660.80
2	<function median at 0x7f5d15042b00>	1	0.00
3	<function median at 0x7f5d15042b00>	2	125510.82
3	<function median at 0x7f5d15042b00>
2	<function median at 0x7f5d15042b00>
4	<function median at 0x7f5d15042b00>	5	0.00
1	<function median at 0x7f5d15042b00>	10	57369.61
...	<function median at 0x7f5d15042b00>	7	0.00
...	<function median at 0x7f5d15042b00>	3	75075.31

```

2
9999 <function median at 0x7f5d15042b00>      4  130142.79
1

```

	HasCrCard	IsActiveMember	EstimatedSalary	Exited
0	1	1	101348.88	1
1	0	1	112542.58	0
2	1	0	113931.57	1
3	0	0	93826.63	0
4	1	1	79084.10	0
...
9995	1	0	96270.64	0
9996	1	1	101699.77	0
9997	0	1	42085.58	1
9998	1	0	92888.52	1
9999	1	0	38190.78	0

```
[10000 rows x 14 columns]
```

Spilt the data into training and testing

```

from sklearn.model_selection import train_test_split
train_size=0.8
X=df.drop(columns=['Age']).copy()
Y=df['Age']
X_train,X_rem,Y_train,Y_rem=train_test_split(X,Y,train_size=0.8)
test_size=0.5
X_valid,X_test,Y_valid,Y_test=train_test_split(X_rem,Y_rem,test_size=0
.5)
print(X_train.shape),print(Y_train.shape)
print(X_valid.shape),print(Y_valid.shape)
print(X_test.shape),print(Y_test.shape)

(8000, 13)
(8000,)
(1000, 13)
(1000,)
(1000, 13)
(1000,)

```